


NDT, Section 7.3.2: on a stepped objective the step size decides the answer, and not monotonically -- 0.05 and 0.5 reach the optimum, 0.1 and 0.3 stall

Eq. (7.16) objective over translation, sliced at the true yaw (8 deg)

step 0.5: 13 iters, score -4.1 | step 0.1: 3 iters, score 1069 (both report converged)



Slice through the optimum. The basin is deep and clean, but the surface steps: points cross voxel boundaries and ndt_score divides by the number matched, so the value jumps (note $x = -0.45$). Finite differences have to straddle those steps to mean anything.

