

1. Project Overview

1.1 Purpose

This repository provides **open-source, simulation-based examples** of the algorithms in *Principles of Indoor Positioning and Indoor Navigation* (Chapters 2–8).

The repo is intended as a **companion to the book**: readers can read the equations and derivations in the text and then run concrete, minimal code examples that directly implement those equations on simulated datasets.

1.2 Target Users

- Master's students in navigation / robotics / geodesy / EE / CS.
- PhD students and early-career researchers entering indoor positioning.
- Engineers who want a **reference implementation** of classical algorithms before building production systems.

1.3 Goals

The repo should:

- Demonstrate key algorithms from Ch.2–8 on **small, reproducible simulations**:
 - **Ch.2:** coordinate systems & attitude.
 - **Ch.3:** LS/WLS, robust LS, KF/EKF/UKF/PF, FGO.
 - **Ch.4–7:** RF positioning, fingerprinting, PDR/sensors, SLAM.
 - **Ch.8:** practical sensor fusion.
- Provide **reproducible experiments** based on open simulation datasets.
- Offer **simple, inspectable reference implementations** rather than production systems.
- Support **per-chapter examples** that map back to the book.

New Goal: Equation-Level Traceability

For every important equation in the book that is implemented in code, there must be a **clear mapping**:

- **Book → Code**
 - For a given equation number (e.g. Eq. (3.12)), users can quickly locate:
 - The implementing function/class.
 - The tests and example notebooks that exercise it.
- **Code → Book**

- From a function or class, users can see:
 - Which equation(s) it implements.
 - The chapter/section context.

Design constraints:

- Mapping must be:
 - **Searchable** via plain text (e.g. searching "Eq. (3.12)" in the repo).
 - **Maintainable** when new code is added.
 - Visible in **docstrings, comments, docs, and notebooks**.

1.4 Out of Scope

- Advanced techniques from Chapter 9 (crowdsourcing, collaborative, deep AI PDR, RIS) beyond small demos.
 - Real-time deployment, mobile apps, or large-scale real-data pipelines.
 - Full SLAM frameworks or high-performance mapping stacks.
-

2. Scope & Use Cases

2.1 Scope

The repo implements **simulation-first** reference versions of:

- Coordinate transforms (body/map/ENU/NED/LLH) and attitude representations.
- LS/WLS and robust LS; KF/EKF/UKF/PF; basic FGO wrappers.
- RF point positioning (TOA / two-way TOA / TDOA / AOA / RSS).
- Fingerprinting (Wi-Fi / magnetic / hybrid; deterministic & probabilistic).
- Proprioceptive sensor models (IMU, wheel odom, PDR, barometer, magnetometer).
- Minimal 2D SLAM & sensor fusion demos (Loosely vs Tightly coupled; observability).

2.2 Typical Use Cases

- “Run a notebook to reproduce a figure similar to chapter X.”
- “Swap KF ↔ FGO on the **same** simulated trajectory.”
- “Compare RF trilateration vs TDOA vs AOA in a toy floor.”
- “Compare k-NN fingerprinting vs a simple ML model on a synthetic RF map.”

This section sets **what problems the repo must solve**, not just what code exists.

3. High-Level Architecture

Repo layout:

```
ipin-examples/
  core/
    coords/
    estimators/
    rf/
    sensors/
    sim/
    eval/
  ch2_coords/
  ch3_estimators/
  ch4_rf_point_positioning/
  ch5_fingerprinting/
  ch6_dead_reckoning/
  ch7_slam/
  ch8_sensor_fusion/
  data/
    sim/
    real/  # optional, small demo logs only
  notebooks/
  docs/
  tools/
  check_equation_index.py
```

Directory roles:

- **core/**

Reusable math & models used across chapters (never import from chX_...):

- coords/ – frames, LLH↔ECEF↔ENU, ENU↔NED, attitude conversions.

- estimators/ – LS/WLS, robust LS, KF/EKF/UKF/PF, FGO wrappers.
 - rf/ – TOA/TDOA/AOA/RSS models, DOP utilities.
 - sensors/ – IMU, wheel odom, PDR, mag, barometer models.
 - sim/ – trajectory generators, scenario definitions, noise injection.
 - eval/ – error metrics, CDFs, NEES/NIS, DOP.
 - **chX_.../**
 Thin chapter-specific examples:
 - Wiring code & scripts.
 - Plots and figure reproduction.
 - Minimal glue code between core/ and data.
 - **data/sim/**
 Standardized **simulation datasets** (Section 5).
 - **data/real/**
 Optional, very small demonstration logs (not the main focus).
 - **notebooks/**
 One notebook (or small set) per chapter, referencing equations and pointing to core/ modules.
 - **docs/**
 Markdown docs per chapter + equation_index.yml and usage docs.
 - **tools/**
 CI/maintenance scripts, especially equation mapping checker.
-

4. Core Functional Requirements (Basic Functions)

For each core submodule, the design doc specifies:

- **Purpose**
- **Key data structures**
- **Required functions / classes** (with rough signatures)
- **Equation mapping** expectations (docstring conventions, index entries)
- **Owner:** navigation engineer vs software engineer.

4.1 core/coords

Purpose

- Implement coordinate and attitude foundations of Ch.2.

Basic functions

- Geodetic / ECEF / local frames:
 - llh_to_ecef(llh) -> np.ndarray
 - ecef_to_llh(ecef) -> np.ndarray
 - ecef_to_enu(ecef, ref_llh) -> np.ndarray
 - enu_to_ecef(enu, ref_llh) -> np.ndarray
 - enu_to_ned(enu) -> np.ndarray
 - ned_to_enu(ned) -> np.ndarray
- Attitude conversions:
 - rpy_to_rotmat(roll, pitch, yaw) -> np.ndarray
 - rotmat_to_rpy(R) -> Tuple[float, float, float]
 - quat_to_rotmat(q) -> np.ndarray
 - rotmat_to_quat(R) -> np.ndarray

Equation mapping

- All core transforms have docstrings like:
 - "Implements Eq. (2.x) in Chapter 2: ..."
- Each implemented equation appears in docs/equation_index.yml.

Unique tasks

- Support multiple map frames (multi-floor, building frames).
- Enforce consistent frame naming/metadata used by RF, PDR, SLAM.

A. Replacement text for Section 4.2 core/estimators

Use this to replace your current “4.2 core/estimators” subsection.

4.2 core/estimators

4.2.1 Purpose

This module implements the core estimation and filtering algorithms from Chapter 3 and exposes them through reusable APIs that other modules (core/rf, core/sensors, chX...) can call.

Goals:

- **Provide reference implementations of:**

- Least squares (LS), weighted LS (WLS), and robust LS.
 - Linear KF, EKF, UKF, PF.
 - Factor graph optimization (FGO) with basic numerical solvers.
 - Maintain equation-level traceability to Chapter 3:
 - Each algorithm implementation references the exact equation numbers it follows.
 - Reuse the same estimator APIs across RF positioning, PDR, SLAM, and fusion examples.
-

4.2.2 Key abstractions & data structures

- State vector
 - State = np.ndarray of shape (n,).
- Covariance matrix
 - Covariance = np.ndarray of shape (n, n).
- Measurement vector
 - Measurement = np.ndarray of shape (m,).
- ProcessModel (for KF/EKF/UKF/FGO)
 - Encapsulates system dynamics:
 - $f(x, u, dt) \rightarrow x_{\text{next}}$: nonlinear state propagation (Eq. (3.21) for EKF, Eq. (3.25) for UKF).
 - $F(x, u, dt) \rightarrow F_k$: Jacobian $\partial f / \partial x$ for EKF (used in Eq. (3.22)).
 - $Q(x, u, dt) \rightarrow Q_k$: process noise covariance.
- MeasurementModel
 - Encapsulates sensor model:
 - $h(x) \rightarrow z_{\text{pred}}$: linear or nonlinear measurement (e.g. Eq. (3.8) for KF, Eq. (3.21) measurement part).
 - $H(x) \rightarrow H_k$: Jacobian $\partial h / \partial x$ for EKF.
 - $R(x) \rightarrow R_k$: measurement noise covariance ($\Sigma_{\{w,z,k\}}$).
- Factor / FactorGraph (for FGO)
 - Factor:

- Holds residual function $r(x_{\text{subset}})$ and its Jacobian (or automatic differentiation stub).
- FactorGraph:
 - Collection of variables and factors representing the MAP optimization problem of Eq. (3.35)–(3.38).

All estimators share these abstractions so that changing the estimator does not require changing the sensor or motion models.

4.2.3 Least Squares (LS / WLS / Robust LS)

Relevant equations (Chapter 3)

- Cost function:
 - $J(x) = \sum (y_i - h_i(x))^2 \rightarrow \text{Eq. (3.1).}$
- Linear model normal equations:
 - $H^T H \hat{x} = H^T y \rightarrow \text{Eq. (3.2).}$
 - $\hat{x} = (H^T H)^{-1} H^T y \rightarrow \text{Eq. (3.3).}$
- Nonlinear LS optimality condition:
 - $\sum (y_i - h_i(x)) \nabla h_i(x) = 0 \rightarrow \text{Eq. (3.4).}$
- Weighted LS cost (no explicit equation number, defined in the text in 3.1.1):
 - $J(x) = \sum w_i (y_i - h_i(x))^2.$

Planned API (in core/estimators/least_squares.py)

- def linear_least_squares(H: np.ndarray, y: np.ndarray) -> np.ndarray:
 - Implements Eq. (3.2)–(3.3).
 - Docstring must include:
 - “Implements Eqs. (3.2)–(3.3) (normal equations and closed-form LS solution).”
- def weighted_least_squares(H: np.ndarray, y: np.ndarray, W: np.ndarray) -> np.ndarray:
 - W is diagonal or full weight (inverse covariance).
 - Generalizes Eq. (3.2)–(3.3) to WLS.
 - Docstring: “Implements weighted least squares as defined in Section 3.1.1.”

- **def gauss_newton_solve(h_fun, jac_fun, y, x0, max_iters, tol) -> np.ndarray:**
 - Implements iterative solution of Eq. (3.4): residual gradient=0 via Gauss–Newton.
 - Used for nonlinear LS in RF / SLAM.
- **def robust_least_squares(...) -> np.ndarray:**
 - Wraps LS with robust loss functions (Huber, Cauchy, etc.) as described in Section 3.1.1.
 - Docstring references Table 3.1 + robust LS text, even if no explicit equation number.

Equation mapping requirements

- **linear_least_squares docstring:** references Eqs. (3.2)–(3.3).
 - **gauss_newton_solve:** references Eq. (3.4).
 - **equation_index.yml must map:**
 - "Eq. (3.2)" and "Eq. (3.3)" → core/estimators/least_squares.py::linear_least_squares.
 - "Eq. (3.4)" → core/estimators/least_squares.py::gauss_newton_solve.
-

4.2.4 Kalman Filter Family (KF / EKF / UKF)

(a) Linear Kalman Filter

Relevant equations (Chapter 3)

- **Linear measurement model:**
 - $z_k = H_k x_k + w_{\{z,k\}} \rightarrow \text{Eq. (3.8).}$
- **Likelihood mean/covariance:**
 - $E[z_k | x_k] = H_k x_k, \text{Cov}[z_k | x_k] = \Sigma_{\{w,z,k\}} \rightarrow \text{Eq. (3.9).}$
- **Propagation model and prior mean/cov:**
 - $x_{k,k-1} = F_k x_{\{k-1\}} + u_k + w_{\{x,k-1\}} \rightarrow \text{Eq. (3.11).}$
 - $P_{\{k,k-1\}} = F_k \Sigma_{\{x,k-1\}} F_k^T + \Sigma_{\{w,u,k\}} \rightarrow \text{Eq. (3.12).}$
- **MAP derivation and posterior:**
 - Posterior density (3.13), MAP via derivative (3.15), closed form (3.16).
- **KF update form:**

- $\hat{x}_{\{k, \text{MAP}\}} = x_{\{k, k-1\}} + K_k (z_k - H_k x_{\{k, k-1\}}) \rightarrow \text{Eq. (3.17).}$
- $K_k = P_{\{k, k-1\}} H_k^T (H_k P_{\{k, k-1\}} H_k^T + \Sigma_{\{w, z, k\}})^{-1} \rightarrow \text{Eq. (3.18).}$
- Covariance: $\Sigma_{\{x, k\}} = P_{\{k, k-1\}} - F_k K_k H_k P_{\{k, k-1\}} \rightarrow \text{Eq. (3.19).}$
- Summary of “five equations” → grouped as Eq. (3.20).

Planned API (in core/estimators/kalman.py)

- class KalmanFilter:
 - `__init__(self, F, Q, H, R)` – constant matrices or callables.
 - `predict(self, u: np.ndarray = None)`
 - Implements propagation from Eqs. (3.11), (3.12) and the first two lines of (3.20).
 - `update(self, z: np.ndarray)`
 - Implements Eqs. (3.17)–(3.19) and full set in (3.20).

Equation mapping requirements

- predict docstring: “Implements the prediction step of the linear Kalman filter (Eqs. (3.11), (3.12), (3.20)).”
 - update docstring: “Implements the update step (Eqs. (3.17)–(3.20)).”
 - equation_index.yml:
 - "Eq. (3.11)", "Eq. (3.12)", "Eq. (3.17)", "Eq. (3.18)", "Eq. (3.19)", "Eq. (3.20)" → KalmanFilter.predict / KalmanFilter.update.
-

(b) Extended Kalman Filter (EKF)

Relevant equations (Chapter 3)

- Nonlinear state & measurement model:
 - $x_k = f(x_{\{k-1\}}, u_k) + w_k, z_k = h(x_k) + v_k \rightarrow \text{Eq. (3.21).}$
- EKF prediction:
 - $\hat{x}_k^- = f(\hat{x}_{\{k-1\}}, u_k),$
 - $P_k^- = F_{\{k-1\}} P_{\{k-1\}} F_{\{k-1\}}^T + Q \rightarrow \text{Eq. (3.22).}$
- EKF update (on the next page, same section; equation number continues, but we primarily tie to Eq. (3.21)–(3.23)).

Planned API

- **class ExtendedKalmanFilter(KalmanFilter):**
 - Accepts ProcessModel and MeasurementModel objects instead of fixed matrices.
 - **predict(self, u, dt):**
 - Uses f , F , and $Q \rightarrow$ Eq. (3.21) and (3.22).
 - **update(self, z):**
 - Uses h , H , R and standard KF structure, referencing EKF update equations in the EKF subsection.

Equation mapping

- predict docstring: “Implements EKF prediction (Eqs. (3.21)–(3.22)).”
 - update docstring: “Implements EKF update (EKF section following Eq. (3.21)).”
 - equation_index.yml: map "Eq. (3.21)", "Eq. (3.22)" to ExtendedKalmanFilter.predict.
-

(c) Unscented Kalman Filter (UKF)

Relevant equations (Chapter 3)

- Sigma-point generation:
 - $\chi_0 = \hat{x}_{k-1}$, $\chi_i = \hat{x}_{k-1} + \delta_i$, $\chi_{i+n} = \hat{x}_{k-1} - \delta_i$, $i=1..n \rightarrow$ Eq. (3.24).
- Sigma-point propagation through f (process) and h (measurement):
 - $\chi_i^- = f(\chi_i, u_k) \rightarrow$ Eq. (3.25).
- UKF update and gain:
 - K_k , \hat{x}_k , P_k defined via cross-covariances and measurement covariances \rightarrow Eq. (3.30).

Planned API

- **class UnscentedKalmanFilter:**
 - **__init__(self, process_model, measurement_model, alpha, beta, kappa)**
 - **predict(self, u, dt)**
 - Generates sigma points (Eq. (3.24)), propagates via f (Eq. (3.25)), computes predicted mean/cov.
 - **update(self, z)**

- Computes predicted measurement, cross-covariance, Kalman gain, and posterior state following Eq. (3.30).

Equation mapping

- Docstrings reference Eqs. (3.24), (3.25), (3.30) explicitly.
 - equation_index.yml maps these equations to the UKF methods.
-

4.2.5 Particle Filter (PF)

Relevant equations (Chapter 3)

- Recursive Bayes update:
 - $p(x_{-k} | z_{1:k}) \propto p(z_k | x_{-k}) p(x_{-k} | z_{1:k-1}) \rightarrow$ Eq. (3.32).
- Sampling step:
 - $x_{-k}^{(i)} \sim p(x_{-k} | x_{\{k-1\}}^{(i)}) \rightarrow$ Eq. (3.33).
- Weight update:
 - $\dot{w}_{-k}^{(i)} = w_{\{k-1\}}^{(i)} p(z_k | x_{-k}^{(i)}) \rightarrow$ Eq. (3.34).

Planned API (in core/estimators/particle.py)

- class ParticleFilter:
 - Attributes: particles ($N \times n$), weights (N), process_model, measurement_model.
 - predict(self, u, dt):
 - Sample new particles using the process model (Eq. (3.33)).
 - update(self, z):
 - Update weights via likelihood (Eq. (3.34)), normalize, resample.
 - estimate(self) -> np.ndarray:
 - Return weighted mean or best-weight particle.

Equation mapping

- predict docstring references Eq. (3.33).
 - update docstring references Eq. (3.32)–(3.34).
 - equation_index.yml: "Eq. (3.32)", "Eq. (3.33)", "Eq. (3.34)" → ParticleFilter.
-

4.2.6 Factor Graph Optimization (FGO) & Numerical Methods

Relevant equations (Chapter 3)

- MAP for full trajectory:
 - $\hat{X}_{\text{MAP}} = \underset{X}{\operatorname{argmax}} p(X | Z) = \underset{X}{\operatorname{argmax}} p(Z | X) p(X) / p(Z) \rightarrow \text{Eq. (3.35).}$
- Simplified MAP form:
 - $\hat{X}_{\text{MAP}} = \underset{X}{\operatorname{argmax}} \ell(X; Z) p(X) \text{ with } \ell(X; Z) \propto p(Z | X) \rightarrow \text{Eqs. (3.36)–(3.37).}$
- Conversion from product of Gaussians to sum of squared residuals (negative log-posterior) → Eq. (3.38) and following.
- Gradient descent update:
 - $x_{k+1} = x_k + \alpha d \rightarrow \text{Eq. (3.42).}$
- Descent condition:
 - $f(x_{k+1}) = f(x_k + \alpha d) < f(x_k) \rightarrow \text{Eq. (3.43).}$

Planned API (in core/estimators/fgo.py and core/estimators/optim.py)

- class Factor:
 - def residual(self, x: np.ndarray) -> np.ndarray:
 - def jacobian(self, x: np.ndarray) -> np.ndarray:
- class FactorGraph:
 - Stores variable ordering, list of factors.
 - def evaluate(self, x: np.ndarray) -> Tuple[r, J]: returns stacked residuals and Jacobian.
- def gradient_descent_step(x, grad, alpha) -> np.ndarray:
 - Encodes Eq. (3.42) and Eq. (3.43) logic.
- def gauss_newton_step(x, J, r) -> np.ndarray:
 - Equivalent to LS step on residuals derived from Eq. (3.38).
- def levenberg_marquardt_step(...) -> np.ndarray:
 - LM step bridging gradient descent and Gauss–Newton, consistent with 3.4.1 discussion.
- def solve_fgo(graph: FactorGraph, x0, method='gn', max_iters=..., tol=...) -> np.ndarray:
 - High-level solver that iteratively applies GN/LM using the factor graph.

Equation mapping

- `solve_fgo` docstring references Eqs. (3.35)–(3.38) (MAP via factor graph).
 - `gradient_descent_step` docstring references Eqs. (3.42)–(3.43).
 - `equation_index.yml` ties these equations to the corresponding functions.
-

4.2.7 Simulation data requirements (for estimators)

Although most estimators are exercised through chapter-specific scenarios, core/estimators also needs simple synthetic datasets for unit tests and minimal examples:

- `toy_ls_linear/`
 - Small linear regression / linear system:
 - H ($m \times n$), y (m), known true x_{true} , noise variance.
 - Used to test:
 - `linear_least_squares`, `weighted_least_squares`, `robust_least_squares`.
 - Must be documented in Section 5.2 as one of the dataset families.
- `toy_kf_1d_cv/`
 - 1D constant-velocity (CV) model:
 - State: [position, velocity].
 - Measurements: noisy position only.
 - Used to test:
 - KalmanFilter convergence vs. analytical solution.
- `toy_nonlinear_1d/`
 - 1D or 2D nonlinear tracking:
 - Nonlinear measurement, e.g. range², to exercise EKF, UKF, PF, and FGO.
 - Used to:
 - Compare KF vs. EKF vs. UKF vs. PF on the same synthetic data.
 - Test FGO smoothing vs. recursive filters.

These datasets should be stored under `data/sim/` (see Section 5.2) and used both by:

- Unit tests in `tests/test_estimators_*.py`.
- A small `ch3_estimators` notebook that reproduces selected plots / behaviors from Chapter 3.

4.3 core/rf

Purpose

- RF signal measurement models from Ch.4–5.

Basic functions

- Ranging / angles:
 - `toa_range(tx_pos, rx_pos, c, clock_bias)`
 - `two_way_toa_range(...)`
 - `tdoa_range_diff(anchor_i, anchor_j, state, c)`
 - `aoa_bearing(anchor_pos, state)`
- RSS:
 - `rss_pathloss(tx_power, distance, n, sigma_shadow)`
- DOP:
 - `compute_dop(geometry, weights=None) -> dict with HDOP/VDOP/PDOP.`

Equation mapping

- Each measurement model documents which equations in Ch.4 it implements (e.g. path-loss, hyperbolic TDOA, AOA geometry).

Unique tasks

- Support multiple technologies (Wi-Fi, BLE, UWB, 5G) through parameterization rather than separate code.

4.4 core/sensors

Purpose

- Proprioceptive & environmental sensor models from Ch.6.

Basic functions

- IMU:
 - Continuous/discrete error models (bias, noise, RW).
 - Simple strapdown integration for 2D/3D.
- PDR:

- Step detection, step length models, heading from IMU/mag.
- Environmental:
 - Barometer altitude, floor detection.
 - Magnetometer heading and magnetic fingerprint stubs.
- Wheel odometry:
 - Differential drive, steering model, noise models.

Equation mapping

- IMU propagation, error models, ZUPT updates etc. point to Ch.6 equations.

Unique tasks

- Define a **unified sensor packet** structure used across simulation and logs.

4.5 core/sim

Purpose

- Shared simulation tools for all examples.

Basic functions

- Trajectories:
 - generate_2d_trajectory(shape, params, dt, duration)
 - generate_3d_trajectory(...)
- Scenarios:
 - Rooms/floor plans, beacons/anchors, obstacles.
- Sensor noise injection:
 - add_imu_noise(traj, model_params)
 - simulate_rf_measurements(traj, anchors, rf_config)

Unique tasks

- Scenario configurations in JSON/YAML so all chapters share the same definitions.
- **4.6 core/eval**
- **Purpose**
 - Evaluation and visualization utilities shared across chapters.
 - Provide a small, consistent API for:
 - computing errors and consistency metrics, and
 - generating publication-quality, vector graphics for trajectories, errors, and algorithm comparisons.

- **Basic functions (numerical evaluation)**
 - Position errors:
 - `compute_position_errors(truth, est)` -> dict
 - Per-epoch error vectors (ENU / NED).
 - `compute_rmse(errors)` -> float | dict
 - Scalar RMSE or per-axis RMSE.
 - `compute_error_stats(errors)` -> dict
 - mean / median / percentiles.
 - Consistency metrics:
 - `compute_nees(truth, est, cov)` -> np.ndarray
 - `compute_nis(innov, S)` -> np.ndarray
 - helpers to check NEES/NIS against χ^2 bounds.
 - RF / geometry utilities:
 - `compute_dop(geometry, weights=None)` -> dict (HDOP/VDOP/PDOP), reusing core/coords for frame handling.
- 574882d9-b20d-4e2c-a4e3-c94a8d9...
- **Basic functions (visualization)**

All plotting helpers live in `core/eval/plots.py`. They must:

 - accept plain NumPy arrays / dicts (no hidden global state), and
 - return a `matplotlib.figure.Figure` (the caller decides whether to show or save).
- Planned APIs:
 - Trajectory / map views
 - `plot_trajectory_2d(truth_xy, est_xy_dict, anchors_xy=None)` -> Figure
 - truth vs. one or more estimates on a 2D floor plan.
 - optional RF anchors / landmarks overlay.
 - `plot_trajectory_3d(truth_xyz, est_xyz_dict, anchors_xyz=None)` -> Figure
 - simple 3D view for SLAM / multi-floor examples.
- Error over time
 - `plot_position_error_time(errors_dict, dt, axes="enu")` -> Figure
 - per-axis error vs. time, multiple algorithms on one plot.
- Error distributions
 - `plot_error_hist(errors_dict, bins=...)` -> Figure
 - `plot_error_cdf(errors_dict)` -> Figure
 - used heavily in Ch.4–5 to compare RF / fingerprinting methods.
- 574882d9-b20d-4e2c-a4e3-c94a8d9...
- - Consistency plots
 - `plot_nees(nees, chi2_bounds)` -> Figure
 - `plot_nis(nis, chi2_bounds)` -> Figure
 - RF geometry / DOP
 - `plot_rf_geometry(anchors_xy, traj_xy=None)` -> Figure
 - `plot_dop_map(dop_grid, floorplan=None)` -> Figure
 - visualize how beacon layouts (or AP layouts) affect HDOP / VDOP.
 - SLAM / map views (minimal)
 - `plot_occupancy_grid(grid, poses=None, landmarks=None)` -> Figure
 - `plot_factor_graph_skeleton(nodes, edges)` -> Figure (optional, diagnostic only).
- **Standard vector export**
- To make results easy to reuse in papers and slides, every plotting helper must work with a single, shared export utility:

- `save_figure(fig, out_dir, name, formats=("svg", "pdf")) -> list[path]`
 - ensures:
 - vector formats by default (.svg, .pdf),
 - deterministic naming, and
 - creation of the output directory if needed.
 - **File and naming conventions**
 - All chapter examples should save figures into:
 - `chX_.../figs/` for scripts, or
 - `notebooks/figs/chX_.../` for notebooks.
 - File names must follow:
 - `ch3_estimators_kf_vs_pf_rmse.svg`
 - `ch4_rf_toa_vs_tdoa_cdf.pdf`
 - `ch6_pdr_drift_correction_trajectory.svg`
 - Notebooks *embed* the figure inline, but also call `save_figure` so users get reusable vector files with a single run.
 - **Visual style guidelines**
 - To keep plots visually consistent across the repo:
 - Always label axes with units, e.g. Position error [m], Time [s], Heading [deg].
 - Use the same color / linestyle mapping for estimator families across chapters:
 - LS/WLS: solid blue
 - KF/EKF/UKF: solid green
 - PF: dashed orange
 - FGO / smoothing: solid red
 - Keep backgrounds white and gridlines light (MATLAB/Matplotlib default is fine).
 - Avoid heavy custom styling in core; chapter examples can adjust if needed.
 - **Unique tasks**
 - Define a small set of “standard plots” per chapter (see Section 7) and implement them using `core/eval/plots.py`.
 - Make sure every example script / notebook:
 - prints a small numerical summary (RMSE, final bias, etc.), and
 - generates at least one vector figure via `save_figure(...)`.
 - Ensure that evaluation APIs are reused across RF, PDR, SLAM, and fusion examples so that algorithm comparisons are meaningful and visually comparable.
-

5. Simulation Datasets & Open Data

5.1 Principles

- All primary examples must run **only on data in data/sim/**.
- Datasets must be:
 - Small enough to clone and run quickly.
 - Reproducible (fixed seeds, recorded configs).
 - Clearly licensed (e.g. CC-BY-4.0) and documented in `docs/data.md`.

5.2 Dataset Families

Planned families under data/sim/:

- rf_2d_floor/
 - 2D floor map, anchor positions, true trajectories, TOA/TDOA/AOA/RSS measurements.
- wifi_fingerprint_grid/
 - Reference points on a grid, RSS vectors, test trajectories.
- pdr_corridor_walk/
 - IMU time series, reference path, step labels, floor labels.
- slam_lidar2d/
 - 2D occupancy grid, range-bearing measurements, ground truth poses.
- toy_ls_linear/
 - Simple linear system for least-squares testing:
 - Files: H.npy, y.npy, x_true.npy, noise_config.json.
 - Used by: core/estimators/least_squares.py unit tests and ch3_estimators examples.
- toy_kf_1d_cv/
 - 1D constant-velocity system:
 - Files: x_true.npy, z.npy, F.npy, H.npy, Q.npy, R.npy, dt.
 - Used to validate KalmanFilter implementation and tuning.
- toy_nonlinear_1d/
 - Nonlinear state/measurement example:
 - State trajectory file, nonlinear measurement config, noise statistics.
 - Used by: EKF, UKF, PF, and FGO examples to illustrate behavior under non-Gaussian/nonlinear conditions.

For each dataset:

- File formats: CSV / NPZ / HDF5.
- Coordinate frames: clearly defined, link to core/coords.
- Ground truth semantics.

5.3 Optional Real Data

- `data/real/` may contain **small** demo logs (not full research datasets) to show real-world quirks.
 - Not a core dependency for running per-chapter examples.
-

6. Equation-to-Code Mapping Design

This section defines how equation mapping is implemented, enforced, and integrated with core/ and chapter modules.

6.1 Canonical Equation ID Format

- Equations are referenced using a **canonical string**:
 - Eq. (C.NN) for Chapter C, Equation NN.
 - If the book uses section-based numbering, adapt to Eq. (C.SS).

Examples

- Chapter 2, equation 5 → "Eq. (2.5)"
- Chapter 3, equation 12 → "Eq. (3.12)"

Rules

- This exact string must appear in:
 - Docstrings.
 - Comments (for line-level mapping).
 - Notebooks (markdown cells).
 - `docs/equation_index.yml`.

That way, a simple repo search for "Eq. (3.12)" finds all relevant code/tests.

6.2 Code Conventions (Docstrings & Comments)

Requirement

Every function/class that directly implements a book equation must mention it in its **docstring**.

Example: LS

```
def linear_least_squares(A: np.ndarray, b: np.ndarray) -> np.ndarray:
```

```
    """
```

Solve the linear least squares problem $x^* = \text{argmin} \|Ax - b\|_2$.

Implements Eq. (3.12) in Chapter 3 of the book:

$$x^* = (A^T A)^{-1} A^T b$$

.....

...

Example: EKF predict

```
def predict(self, u: np.ndarray) -> None:
```

.....

Extended Kalman filter prediction step.

Implements Eqs. (3.25)–(3.27) in Chapter 3:

- State propagation: Eq. (3.25)
- Jacobian: Eq. (3.26)
- Covariance propagation: Eq. (3.27)

.....

...

Line-level comments

If a specific line encodes an equation:

```
# Eq. (3.27): P_{k|k-1} = F_k P_{k-1|k-1} F_k^T + Q_k
```

```
P_pred = F @ P @ F.T + Q
```

Design rule

- Whole function = one equation (or block) → docstring reference.
- Single expression or key step = equation → inline comment with the exact Eq. (C.NN) reference.

6.3 Central Equation Index File

Create docs/equation_index.yml (or .json).

Purpose

From the book side, you can look up "Eq. (3.12)" and see:

- Which modules/objects implement it.
- Which tests/notebooks exercise it.

YAML structure (example)

- eq: "Eq. (2.3)"

chapter: 2

description: "ECEF to ENU coordinate transformation"

files:

- path: "core/coords/frames.py"

object: "ecef_to_enu"

- path: "ch2_coords/examples/coord_demo.py"

object: "demo_ecef_to_enu"

tests:

- "tests/test_coords.py::test_ecef_enu_roundtrip"

- eq: "Eq. (3.12)"

chapter: 3

description: "Linear least squares closed form"

files:

- path: "core/estimators/least_squares.py"

object: "linear_least_squares"

notebooks:

- "notebooks/ch3_estimators/ls_vs_ekf.ipynb"

Design choices

- One entry per implemented equation.
- Keys:
 - eq – canonical ID ("Eq. (3.12)").
 - chapter
 - description
 - files – list of {path, object}.
 - Optional: tests, notebooks, notes.

6.4 Notebooks & Docs Conventions

Notebooks

At the top of each chapter notebook, add an “Equations used” list, e.g.:

Equations in this notebook

- Eq. (3.12): Linear least squares solution.
- Eqs. (3.25)–(3.27): EKF prediction equations.

And reference equations in text:

We now implement the linear least squares estimator from Eq. (3.12).

Docs

In docs/ch3_estimators.md:

```
## Equation map
```

- Eq. (3.12): `core/estimators/least_squares.py::linear_least_squares`
- Eqs. (3.25)–(3.27): `core/estimators/kalman.py::ExtendedKalmanFilter.predict`

This gives users three navigation routes:

1. Search "Eq. (3.12)" in the repo.
2. Check docs/equation_index.yml.
3. Read chapter docs listing mappings.

6.5 Tooling: Equation Index Checker Script

Add tools/check_equation_index.py:

Responsibilities:

- Parse docs/equation_index.yml.
- For every {path, object} entry:
 - Verify the module and object exist.
 - Verify the docstring or comments contain the mapped "Eq. (C.NN)".

Sketch (already provided):

- Use importlib to load modules.
- Inspect obj.__doc__.
- Fail CI if any mismatch.

Integrate in CI:

- Run pytest.
- Run python tools/check_equation_index.py.

6.6 Integration into Epics / Phases

For each epic, “definition of done” includes:

- Code & tests written.
- Example notebook added.
- **Equation mapping completed:**
 - Docstring references added.
 - equation_index.yml updated.
 - Notebook/docs list relevant equations.

Concretely:

- **Epic 1 – Coordinates (Ch.2)**
 - All transforms tagged with Eq. (2.x) where applicable.
 - **Epic 2 – Estimators (Ch.3)**
 - LS, KF, EKF, UKF, PF, FGO all mapped to Ch.3 equations.
 - **Epic 3 – RF positioning (Ch.4)**
 - TOA/TDOA/AOA/RSS and DOP models mapped.
 - Subsequent epics (fingerprinting, PDR, SLAM, fusion) follow the same pattern.
-

7. Per-Chapter Example Modules (Unique Tasks)

For each chX_... folder, list:

- Example scripts / notebooks.
- Which core/ functions they exercise.
- Which equations they highlight.

Standard visualization outputs (all chapters)

For each chX_... module, at least one example must:

- call core/eval metrics to compute errors / statistics, and
- call core/eval/plots.py to generate:
 - a trajectory or map view, and
 - at least one error/time or error/CDF plot.

All examples should save figures via `save_figure(...)` into `chX_.../figs/` (scripts) or `notebooks/figs/chX_.../` (notebooks) so that:

- **readers can quickly compare different algorithms or datasets by opening the SVG/PDF files, and**
- **instructors can directly reuse the vector figures in slides and course materials without manual re-plotting.**

7.1 ch2_coords/

Examples

- `demo_frames.py`
 - Build body/map/ENU/NED frames and print transforms.
- Simple plotting script to reproduce or echo Ch.2 figures.

Dependencies

- `core.coords`, `core.sim`, `core.eval`.

7.2 ch3_estimators/

Examples

- `ls_vs_robust_ls.py`
 - Outlier demo: LS vs Huber vs Cauchy.
- `kf_vs_ekf_vs_ukf_vs_pf_vs_fgo.ipynb`
 - 1D/2D toy system comparing estimator families.

Unique tasks

- Compare convergence, RMSE, computation time.
- Show qualitatively how different estimators behave under same conditions.

7.3 ch4_rf_point_positioning/

Examples

- TOA trilateration (LS/WLS).
- TDOA Chan & Fang algorithms.
- AOA OVE / PLE examples.

Dependencies

- `core.coords`, `core.rf`, `core.estimators`, `core.eval`.

7.4 ch5_fingerprinting/

Examples

- Synthetic grid map:
 - Generate RSS fingerprints from core.rf and core.sim.
- Deterministic methods:
 - k-NN, weighted k-NN.
- Probabilistic:
 - Naive Bayes / simple Gaussian likelihood.
- Optional:
 - Small MLP (e.g. 2–3 layers) as a “modern” baseline.

7.5 ch6_dead_reckoning/

Examples

- Foot-mounted PDR with ZUPT:
 - Walk along corridor, accumulate drift, correct with constraints.
- Vehicle odom + IMU:
 - Simple car-like motion on a map.

7.6 ch7_slam/

Examples

- Minimal 2D LiDAR / range-bearing SLAM:
 - Use FGO with pose + landmark factors.
- Loop closure toy example.

7.7 ch8_sensor_fusion/

Examples

- Loosely vs tightly coupled fusion examples:
 - GNSS+IMU style but adapted to indoor sensors.
- Observability demo:
 - Show drift/unobservability vs properly fused sensors.
- Calibration:
 - Simple intrinsic/extrinsic calibration toy.

For each chapter section, the design doc should separate:

- **Core functions reused** from core/.

- Chapter-specific “unique tasks” and plots.
-

8. Non-Functional Requirements

8.1 Language & Tooling

- Python 3.x.
- NumPy, SciPy, matplotlib, optional JAX/PyTorch for ML examples (kept light).
- Packaging via `pyproject.toml` / `poetry` or similar.

8.2 Quality

- Unit tests for all core functions (`tests/`).
- pytest-based CI.
- Equation index checker included in CI.

8.3 Performance

- All examples should run on a typical laptop in **minutes**, not hours.
- Datasets and notebooks: no GPU dependency required.

8.4 Documentation

- `docs/`:
 - Per-chapter docs (overview, how to run examples).
 - `docs/equation_index.yml` and usage instructions.
- `README.md`:
 - Explain structure, target audience, and how to map book ↔ repo.

9. Implementation Roadmap

Use epics as a roadmap for you + contributing engineers.

9.1 Epic 0 – Repo & Infrastructure

- Initialize repo structure.
- Set up:
 - `core/`, `chX_.../`, `data/`, `notebooks/`, `docs/`, `tools/`.
 - CI with pytest + equation checker.

9.2 Epic 1 – Coordinates (Ch.2)

- Implement core/coords.
- Add tests for:
 - LLH \leftrightarrow ECEF \leftrightarrow ENU round-trip.
 - ENU \leftrightarrow NED conversions.
- Implement ch2_coords examples.
- Add Ch.2 equations to equation_index.yml.

9.3 Epic 2 – Estimators (Ch.3)

- Implement LS/WLS + robust LS in core/estimators.
- Implement KF/EKF/UKF/PF skeletons.
- Add simple FGO wrapper.
- Unit tests (simple linear systems).
- ch3_estimators example notebook.
- Update equation_index.yml for Ch.3.

9.4 Epic 3 – RF Positioning (Ch.4)

- Implement measurement models and DOP in core/rf.
- Build RF simulation tools in core/sim.
- ch4_rf_point_positioning examples:
 - TOA, TDOA (Chan/Fang), AOA demos.
- Add Ch.4 equations to equation_index.yml.

9.5 Later Epics – Fingerprinting, Sensors, SLAM, Fusion

- Similar pattern:
 - Spec → core implementation → tests → examples → docs → equation mapping.

10. Working Effectively: SW vs Navigation Engineers

Throughout the design doc, separate:

- **Algorithm / modeling decisions** (navigation engineer):
 - Which equation(s) in the book to implement.
 - What approximations and parameters to use.

- **Implementation details** (software engineer):
 - File layout, function signatures, performance, tests, CI.

The equation mapping system bridges the two:

- Nav engineers can **check correctness** by tracing Eq. → code.
- SW engineers can **preserve traceability** by following docstring + index conventions and the equation checker.